

A Truth Serum for Non-Bayesians: Correcting Proper Scoring Rules for Risk Attitudes*

THEO OFFERMAN

University of Amsterdam

JOEP SONNEMANS

University of Amsterdam

GIJS VAN DE KUILEN

Tilburg University

and

PETER P. WAKKER

Erasmus University

First version received October 2007; final version accepted December 2008 (Eds.)

Proper scoring rules provide convenient and highly efficient tools for incentive-compatible elicitation of subjective beliefs. As traditionally used, however, they are valid only under expected value maximization. This paper shows how they can be generalized to modern (“non-expected utility”) theories of risk and ambiguity, yielding mutual benefits: users of scoring rules can benefit from the empirical realism of non-expected utility, and analysts of ambiguity attitudes can benefit from efficient measurements using proper scoring rules. An experiment demonstrates the feasibility of our generalization.

1. INTRODUCTION

An important problem in mechanism design concerns the elicitation of private information. A design is incentive compatible if the actions of agents, motivated solely by self-interest, nevertheless reveal their true private information (Hurwicz, 1960). A social planner can then use all relevant information to devise the most efficient social allocation. This paper considers the

*A preliminary version of this paper circulated with the title “Is the Quadratic Scoring Rule Really Incentive Compatible?”.

case where the private information concerns subjective beliefs about the likelihood of uncertain events, often modelled using subjective probabilities. This case arises, for instance, when principals rely on the judgements of specialized agents. In the absence of proper incentives, agents may pretend to be more confident about their judgement than they really are, and may not update their beliefs sufficiently, so as to suggest greater ability than they really have (Li, 2007). Manski (2004) presented an historical survey of belief measurement and gave many economic applications.

For belief measurement, incentive-compatible mechanisms had been known at an early stage in the form of proper scoring rules (Brier, 1950; Good, 1952). These scoring rules are particularly efficient mechanisms for eliciting subjective beliefs in an incentive-compatible manner. They use cleverly designed optimization problems where the observation of one single choice suffices to determine the exact quantitative degree of belief of an agent in an uncertain event. Hence, they have recently become popular in experimental economics and game theory (Nyarko and Schotter, 2002). Proper scoring rules have been used in many other fields in the social sciences, including accounting (Wright, 1988), Bayesian statistics (Savage, 1971), business (Staël von Holstein, 1972), education (Echternacht, 1972), finance (Shiller, Kon-Ya and Tsutsui, 1996; Johnstone, 2007*a,b*), medicine (Spiegelhalter, 1986), meteorology (Yates, 1990; Palmer and Hagedorn, 2006), politics (Tetlock, 2005), psychology (McClelland and Bolger, 1994) and other fields (Hanson, 2002; Prelec, 2004).

An alternative way to elicit beliefs that has recently become popular concerns prediction markets on the internet (Wolfers and Zitzewitz, 2004). Here, people trade event-contingent payments regarding uncertain events, such as a guarantee to receive €100 if a Democrat becomes the next President of the United States. If this guarantee is now traded at a price P , then $P/100$ is taken as the market probability of the event. This inference assumes expected value. Johnstone (2007*a*) explained that a financial market can, for many purposes, be analysed as if it were a rational individual, and discussed the role of proper scoring rules in such settings. In a market with agents who maximize the logarithmic utility of wealth, good forecasters are better identified by their performance in terms of proper scoring rules than in terms of their average earnings (Johnstone, 2007*b*).

Whereas all applications of proper scoring rules that we are aware of assume expected value maximization ("risk neutrality"), many deviations have been observed empirically. Under expected utility, risk aversion is the common finding (Bernoulli, 1738). Johnstone (2007*a*) and Winkler and Murphy (1970) discussed the implications of risk aversion for proper scoring rules. Furthermore, many deviations from expected utility have been found, both when probabilities exist ("risk": Allais, 1953; Kahneman and Tversky, 1979) and when probabilities cannot even be specified ("ambiguity": Keynes, 1921; Knight, 1921; Ellsberg, 1961).

This paper extends proper scoring rules from the expected value model as assumed in the 1950s, when proper scoring rules were introduced, to the current state of the art in decision theory. Thus, we can, on the one hand, improve the validity of belief measurement using proper scoring rules. On the other hand, we can use proper scoring rules to obtain more efficient methods for measuring risk and ambiguity attitudes. In economics, probabilities are usually not known, and the importance of quantitative measurements of ambiguity attitudes has been widely understood (Gilboa, 2004; Greenspan, 2004). We show how subjective beliefs and ambiguity attitudes can be isolated from risk attitude in a surprisingly easy way by means of proper scoring rules. Thus, we can correct measurements of subjective beliefs and ambiguity attitudes for non-neutral risk attitudes.

We illustrate the feasibility of our method by an experiment where we measure the subjective beliefs of participants about the future performance of stocks after the provision of information about past performance. The empirical findings confirm the usefulness of our

method. However, violations of additivity of subjective beliefs are reduced but not eliminated by our corrections. Thus, the classical measurements will contain violations of additivity that are partly due to the incorrect assumption of expected value maximization, but partly they are genuine. Subjective beliefs are genuinely non-additive.

The analysis in this paper consists of three parts. The first part (Sections 2–4) considers various modern theories of risk and ambiguity and derives implications for proper scoring rules. The second part (Sections 5 and 6) applies the revealed preference technique to the results of the first part. That is, we do not assume theoretical models to derive implications for empirical observations, but we use empirical observations to derive implications for theoretical models. Section 5 presents the main result of this paper, showing how subjective beliefs can easily be derived from observed choices using what are called risk corrections. Section 6 gives an example to illustrate such a derivation at the individual level. Readers who are only interested in applying our method empirically can skip most of Sections 3–5, only reading Corollary 4. A short direct proof of this result, showing that it holds for general proper scoring rules, is given following the corollary.

The third part of the paper (Sections 7–11) implements our correction method in an experiment. In order to demonstrate the applicability of our method, we use it to investigate some properties of non-additive beliefs and of different implementations of real incentives. Section 7 contains methodological details. Section 8 presents the results regarding the biases that we correct for, and Section 9 explains some implications of the corrections of such biases. Section 10 provides an additional control treatment. The experimental results are discussed in Section 11. A general discussion and conclusions are in Sections 12 and 13. Appendix A makes some technical remarks, Appendix B presents the proofs and Appendix C surveys the implications of modern decision theories for our measurements. The experimental instructions are in Appendix D.

2. PROPER SCORING RULES; DEFINITIONS

Let E denote an event such that an agent is uncertain about whether or not the event obtains, such as whether a stock's value will decrease during the next 6 months. The degree of uncertainty of the agent about E will obviously depend on the information that the agent possesses about E . For most uncertain events, no objective probabilities of occurrence are known, and decisions have to be based on subjective likelihood assessments.

Prospects refer to event-contingent payments. We use the general notation $x_E y$ for a prospect that yields outcome x if event E obtains and outcome y if E^c obtains, with E^c the complementary event not- E . Outcomes are money amounts. Risk concerns the special case of known probabilities. Then, for a prospect $x_E y$, the probability p of event E is known. We identify this prospect with a probability distribution $x_p y$ over money, yielding x with probability p and y with probability $1 - p$.

This paper considers the *quadratic scoring rule* (QSR), the most commonly used proper scoring rule (McKelvey and Page, 1990; Nyarko and Schotter, 2002; Palfrey and Wang, 2007). A QSR prospect

$$(1 - (1 - r)^2)_E (1 - r^2) \quad (1)$$

is offered to the agent, where $0 \leq r \leq 1$ is chosen at the agent's discretion. This number r is a function of E , sometimes denoted r_E , and is called the (uncorrected) reported probability of E . The reasons for using this term will be explained later. If event E has a (objective or subjective) probability p , then, according to all theories considered, r_E will depend only on p , so that we can write it as a function $R(p)$.

Instead of equation (1), we could have used more general prospects $(a - b(1 - r)^2)_E (a - br^2)$ for any $b > 0$ and $a \in \mathbb{R}$. For simplicity, we restrict our attention to $a = b = 1$ as in

equation (1). No negative payments can occur then, so that the agent never loses money. Under the event that happens, the QSR in fact yields 1 minus the squared distance between the reported probability of a clairvoyant (who assigns probability 1 to the event that happens) and the reported probability of the agent (this probability is r under E , and $1 - r$ under E^c). The following observation about a symmetry of the QSR will be useful.

Observation 1. *The QSR for event E presents the same choice of prospects as the QSR for event E^c , with each prospect resulting from r as the reported probability of E identical to the prospect resulting from $1 - r$ as the reported probability of E^c .*

Because of Observation 1, we have

$$r_{E^c} = 1 - r_E \quad (2)$$

and

$$R(1 - p) = 1 - R(p). \quad (3)$$

Hence, we will state many results only for $r \geq 0.5$. The case $r < 0.5$ then follows from equations (2) and (3) applied to E^c .

3. A THEORETICAL ANALYSIS OF PROPER SCORING RULES

In this section, we consider modern decision models for decision making under uncertainty and derive implications for proper scoring rules. As explained in detail in Appendix C, virtually all currently existing models, including multiple priors (Gilboa and Schmeidler, 1989) and Choquet expected utility (Gilboa, 1987; Schmeidler, 1989), evaluate the QSR prospect of equation (1) using the following equation.

$$\text{For } r \geq 0.5: W(E)U(1 - (1 - r)^2) + (1 - W(E))U(1 - r^2). \quad (4)$$

Comments on the case $r < 0.5$ follow later. U is the *utility function*, assumed to be continuous and strictly increasing, and scaled such that $U(0) = 0$. We present a number of cases for W , with each case generalizing the preceding one. Cases 1 and 2 are well known.

Case 1 [Expected value]. U is the identity function, and W is a probability measure P .

Case 2 [Expected utility]. W is a probability measure P .

Case 3 [Probabilistic sophistication (with non-expected utility)]. There exist a probability measure P and a continuous strictly increasing function w , the *probability weighting function*, such that $W(\cdot) = w(P(\cdot))$, $w(0) = 0$, and $w(1) = 1$.

Case 4 [General model]. W satisfies: (i) $W(\emptyset) = 0$; (ii) $W = 1$ for the universal event; (iii) $C \supset D$ implies $W(C) \geq W(D)$.

We distinguish two subcases for Case 3 and, hence, also for Cases 2 and 1.

Subcase a. [Objective probabilities]. The probability measure P is objective, based on statistical data that everyone agrees on.

Subcase b. [Subjective probabilities]. The probability measure P may be subjective and can be revealed from preferences.¹

1. In this paper, the term *subjective probability* is used only for probability judgements that are Bayesian in the sense that they satisfy the laws of probability. In the literature, the term subjective probability has sometimes been used for judgements that deviate from the laws of probability, including cases where these judgements are non-linear transformations of objective probabilities when the latter are given. We use the term (probability) weights or beliefs, depending on the way of generalization, to refer to the latter use.

De Finetti (1937), Savage (1954) and Machina and Schmeidler (1992) gave preference foundations for Cases 1b, 2b and 3b. Case 3 is an interesting intermediate case, with the Bayesian principles violated at the level of decisions but not at the level of beliefs. In the general Case 4, the Bayesian principles are also violated at the level of beliefs. The well-known Allais (1953) paradox shows that expected utility is often violated, so that w and W are non-additive, and we cannot restrict attention to the classical Cases 1 and 2. The well-known Ellsberg (1961) paradox, discussed in detail later, shows that probabilistic sophistication is often violated, so that the general Case 4 has to be considered.

For the general model, the equation to evaluate the QSR prospect of equation (1) for $r < 0.5$ follows from Observation 1:

$$\text{For } r < 0.5 : (1 - W(E^c))U(1 - (1 - r)^2) + W(E^c)U(1 - r^2). \quad (5)$$

For expected value and expected utility, equation (5) agrees with equation (4) and the two equations can be used interchangeably, but for probabilistic sophistication and the general model, equation (5) can be different. The latter separate, “rank-dependent”, way of weighting the outcomes, with weights always summing to 1, was discovered independently by Quiggin (1982) for the special case of risk with given probabilities, and by Schmeidler (1989; first version 1982) for the general model. This idea was the key to the development of the modern non-expected utility theories. It was incorporated in the new version of prospect theory (Tversky and Kahneman, 1992).

Objective probabilities can best be interpreted as a special limiting case of subjective probabilities, a point formalized by Machina (2004). The hypothetical situation of an agent using a subjective probability different from an objective probability, if the latter is given, cannot arise under plausible assumptions (Wakker, 2009).

We now analyse which optimal values r_E are predicted under the various cases considered.

Theorem 1. *In the general model, the optimal choice r in equation (1) satisfies:*

$$\text{If } r > 0.5, \text{ then } r = r_E = \frac{W(E)}{W(E) + (1 - W(E)) \frac{U'(1 - r^2)}{U'(1 - (1 - r)^2)}}. \quad (6)$$

The optimality result for $r < 0.5$ follows from Observation 1 applied to E^c . If $r = 0.5$ is optimal, then it can be a boundary solution for which equation (6) need not hold. We will discuss this case later. Theorem 1 generalizes results, obtained by Winkler and Murphy (1970) for expected utility, to general non-expected utility. The following corollary, first found by Brier (1950), is highly appealing, and is, to the best of our knowledge, the first incentive-compatible result provided in the literature.

Corollary 1. *Under expected value, equation (6) holds for all r and $r = r_E = P(E)$.*

Thus, under expected value, it is in the agent's best interest to report true subjective probabilities. The following section illustrates the extent to which reported probabilities, still commonly equated with subjective probabilities in virtually all applications today, can deviate from subjective probabilities because of empirical deviations from expected value. We next consider the case $r = 0.5$ under expected utility.

Observation 2. Under expected utility with probability measure P , $r_E = 0.5$ implies $P(E) = 0.5$. Conversely, $P(E) = 0.5$ implies $r_E = 0.5$ if risk aversion holds. Under risk seeking, $r_E \neq 0.5$ is possible if $P(E) = 0.5$.

4. DISCREPANCIES BETWEEN SUBJECTIVE PROBABILITIES AND PROPER SCORING RULES: NUMERICAL EXAMPLES

The solutions r presented in this section can be verified by substitution in the implicit equation (6). We will later provide explicit expressions for $R^{-1}(p)$, which we used to find the solutions and to draw Figure 1. We consider two urns each containing 100 balls that are crimson, green, silver or yellow. Urn K ("known") contains 25 balls of each colour, and urn A ("ambiguous") contains the balls in an unknown proportion. One ball will be drawn at random from each urn. C denotes the event of a crimson ball drawn from urn K, with G, S and Y defined similarly. E is the event that the ball drawn from K is not crimson, that is, it is the event $C^c = \{G, S, Y\}$. C_a denotes the event of a crimson ball drawn from urn A, with G_a, S_a and Y_a defined similarly, and $E_a = C_a^c$. Subjects are asked to report their belief in event E

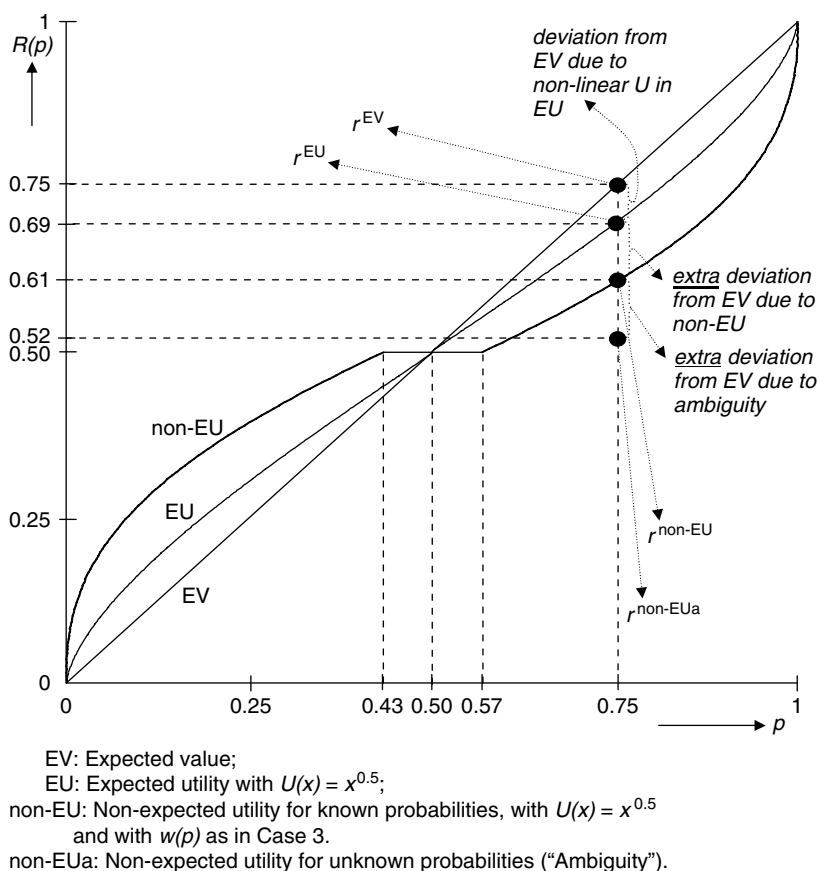


FIGURE 1

Reported probability $R(p)$ as a function of probability p

and are rewarded with a QSR (equation 1). We now consider the four cases presented in the preceding section.

Case 1 [Expected value]. Expected value holds for urn K. Then $r_{EK} = R(0.75) = 0.75$ is optimal in equation (1). The point r_E is depicted as r^{EV} in Figure 1, at $p = 0.75$. Corollary 1 implies that $r_G = r_S = r_Y = 0.25$. The reported probabilities satisfy additivity: $r_G + r_S + r_Y = r_E$. \square

Case 2 [Expected utility]. Expected utility holds for urn K, with $U(x) = x^{0.5}$. We obtain $r_E = R(0.75) = 0.69$, depicted as r^{EU} in Figure 1, at $p = 0.75$. The expected value of the resulting QSR prospect is 0.0031 (i.e. 0.8125–0.8094) less than it was in Case 1. This difference can be interpreted as a risk premium, designating a profit margin for an insurance company. Using equation (2), $r_C = 0.31$, and by symmetry $r_G = r_S = r_Y = 0.31$ too. The reported probabilities violate additivity, with $r_G + r_S + r_Y = 0.93 > 0.69 = r_E$. Because of this violation, the data can directly reveal that expected value, the common assumption in applications of proper scoring rules, does not hold. \square

Case 3 [Non-expected utility with probabilistic sophistication]. Probabilistic sophistication holds for urn K, with $U(x) = x^{0.5}$, and

$$w(p) = (\exp(-(-\ln(p))^\alpha)) \quad (7)$$

with parameter $\alpha = 0.65$ (Prelec, 1998). This function agrees with the prevailing empirical findings (Tversky and Kahneman, 1992; Gonzalez and Wu, 1999; Abdellaoui, 2000; Bleichrodt and Pinto, 2000). We obtain $r_E = R(0.75) = 0.61$, depicted as $r^{\text{non-EU}}$ in Figure 1 at $p = 0.75$. The extra expected-value loss (and, hence, the extra risk premium) relative to Case 2 is 0.0174 (i.e. 0.8094–0.7920). Using equation (2), $r_C = 0.39$, and, by symmetry, $r_G = r_S = r_Y = 0.39$ too. The reported probabilities strongly violate additivity because $r_G + r_S + r_Y = 1.17 > 0.61 = r_E$. \square

The following case describes the most fundamental deviation from expected value and expected utility, driven by ambiguity, a central topic in decision theory today. The case concerns a version of Ellsberg's (1961) paradox.

Case 4 [General case; violation of probabilistic sophistication]. We assume probabilistic sophistication for urn K but consider, in addition, urn A using the general model. If probabilities were assigned to drawings from urn A and probabilistic sophistication were also to hold for this urn, then, in view of symmetry, we should have $P(C_a) = P(G_a) = P(S_a) = P(Y_a)$. Then these probabilities would be 0.25. $P(E_a)$ would then be 0.75, as was $P(E)$ in Case 3. Under probabilistic sophistication combined with non-expected utility as in Case 3, r_{E_a} would be the same as r_E in Case 3 for the known urn, that is, $r_{E_a} = 0.61$. It implies that people would be indifferent between $x_{E,y}$ and $x_{E_a,y}$ for all x and y . The latter condition is, however, typically violated empirically. People usually have a strict preference for known probabilities, implying for instance²

$$100_{E0} > 100_{E_a0}. \quad (8)$$

Consequently, it is impossible to model beliefs about uncertain events E_a with probabilities, and probabilistic sophistication fails. Equation (8) implies that $W(E_a) < W(E)$. Using equation (6), $r_{E_a} < r_E$.³ Given the strong aversion to unknown probabilities that is often found empirically (Camerer and Weber, 1992), we will assume that $r_{E_a} = 0.52$. It is depicted as $r^{\text{non-EU}_a}$ in Figure 1. The extra expected-value loss relative to Case 3 is

2. This also holds if people can choose the three colours to bet on in the ambiguous urn, so that there is no reason to suspect unfavourable compositions.

3. It is easiest to see in equation (6) that $1/r_E$ is decreasing in $W(E)$.

$0.7920 - 0.7596 = 0.0324$. This amount can be interpreted as the ambiguity-premium. Using equation (2), $r_C = 0.48$, and by symmetry $r_G = r_S = r_Y = 0.48$ too. The reported probabilities violate additivity to an extreme degree, with $r_G + r_S + r_Y = 1.44 > 0.52 = r_{E_a}$. \square

Figure 1 illustrates the extent to which reported probabilities can deviate from subjective probabilities, because of violations of expected value. The cases presented in this section concerned $p = 0.75$, but Figure 1 deals with all probabilities p under probabilistic sophistication. In the general model, there are only events and no probabilities, so that the latter cannot be put on the x -axis and no graph can be drawn. For expected utility, a similar figure is in Winkler and Murphy (1970, figure 3). Its pattern was confirmed empirically by Huck and Weizsäcker (2002). The figure illustrates the errors generated by the assumption of expected value maximization from the perspective of modern views in decision theory. Johnstone (2007a, p. 164) gave similar results from the perspective of a mean-variance model.

5. REVEALED PREFERENCE TECHNIQUES TO ELICIT SUBJECTIVE BELIEFS FROM PROPER SCORING RULES

In the preceding sections, we assumed theoretical decision models and derived predictions about reported probabilities in proper scoring rules. This section presents the usual revealed preference technique. That is, we assume that we observe reported probabilities, and we then investigate what we can infer about decision models and their parameters. In particular, we will be interested in inferring subjective probabilities and their generalizations from proper scoring rules.

If we could observe enough general decisions under risk without proper scoring rules, with enough events available with known probabilities (such as those referring to urn K in Case 4), then we could, in principle, reveal the whole function w . Similarly, we could reveal the whole function W if we could observe enough decisions under uncertainty. Then, we could obtain the following concept, which will be central in this paper.

$$B(E) = w^{-1}(W(E)). \quad (9)$$

In general, B assigns value 0 to the vacuous event \emptyset and value 1 to the universal event, and B is increasing in the sense that $C \supset D$ implies $B(C) \geq B(D)$. These properties similarly hold for the composition $W(\cdot) = w(B(\cdot))$, as we saw before. Under probabilistic sophistication (including expected utility and expected value), $B(E)$ agrees with the probability $P(E)$. In all cases in Section 4 up to Case 3, it is indeed the case that $B(E) = 0.75 = P(E)$. Thus, $B(E)$ is a better candidate for measuring subjective beliefs than r_E , the value still commonly used in applications of proper scoring rules today. Whenever subjective probabilities exist, B measures them correctly, irrespective of the risk attitude. B is what results from r_E after correction for non-neutral risk attitudes. We call B the *(risk-)corrected reported probability*.

Case 4 showed that decisions sometimes cannot be modelled using subjective probabilities. In particular, B in equation (9) will not be a probability measure in Case 4. Yet we think that B is a better candidate to reflect subjective beliefs of likelihood than the uncorrected reported probabilities. Risk attitude is a behavioural component rather than a component reflecting beliefs and, hence, it should be filtered out from belief assessments. Many studies of direct judgements of belief have supported the thesis that subjective beliefs cannot be modelled through probabilities (Shafer, 1976; McClelland and Bolger, 1994; Tversky and Koehler, 1994), so that B will violate additivity. Bounded rationality is an extra reason to expect that subjective beliefs will violate the laws of probability (Aragones *et al.*, 2005; Charness and Levin, 2005).

Example 1. Consider Case 4. The belief component $B(E_a)$ is estimated to be $w^{-1}(W(E_a)) = w^{-1}(0.52) = 0.62$. This value implies that B must violate additivity. Under additivity, we would have $B(C_a) = 1 - B(E_a) = 0.38$, and then, by symmetry, $B(G_a) = B(S_a) = B(Y_a) = 0.38$, so that $B(G_a) + B(S_a) + B(Y_a) = 3 \times 0.38 = 1.14$. This value should then equal $B\{G_a, S_a, Y_a\} = B(E_a) = 0.62$, but it does not. Additivity is violated and B is no probability measure.

Of the total deviation of $r_{E_a} = 0.52$ from 0.75, which is 0.23, a part of $0.06 + 0.08 = 0.14$ is the result of deviations from risk neutrality that distorted the measurement of $B(E_a)$. The remaining 0.09 is not a distortion in the measurement of belief. It rather shows that belief is genuinely non-additive. \square

The measurement of B through entire measurements of w and W is laborious, in particular because of interactions with utility (Tversky and Kahneman, 1992, p. 311; Abdellaoui, Vossman and Weber, 2005). The following results prepare for a tractable measurement of B . Whereas the expression of r in terms of W in Theorem 1 was implicit, we now present an explicit expression of its inverse, that is, of W in terms of r . For easy later reference, we state the result for $B = w^{-1}(W)$ instead of W .

Corollary 2. *For the optimal choice $r = r_E$:*

$$\text{If } r > 0.5, \text{ then } B(E) = w^{-1} \left(\frac{r}{r + (1-r) \frac{U'(1-(1-r)^2)}{U'(1-r^2)}} \right). \quad (10)$$

We next display the special case of $W = w(P)$ (with P objective or subjective), in which case $B = P = R^{-1}(r)$.

Corollary 3. *Under probabilistic sophistication, we have for the optimal choice $r = R(p)$:*

$$\text{If } r > 0.5, \text{ then } p = w^{-1} \left(\frac{r}{r + (1-r) \frac{U'(1-(1-r)^2)}{U'(1-r^2)}} \right). \quad (11)$$

As it so happens, the right-hand sides in equations (10) and (11) are identical. This allows a particularly convenient way to measure B .

Corollary 4. *Assume the general model (equations 4 and 5). For an event E with $r_E = r$, we can find the objective probability p with the same value $R(p) = r$, and then we can conclude that $B(E) = p$. That is,*

$$\text{If } r_E = r > 0.5, \text{ then } B(E) = R^{-1}(r). \quad (12)$$

The corollary is useful for empirical applications because all the terms involved are easily observable. The corollary is the only implication of our theoretical analysis that is needed for our application. Two points underlie the corollary. First, it can be applied where proper scoring rules uniquely identify the underlying subjective factors so that the inverse function R^{-1} can be defined. Second, $R(p)$ and r_E optimize the same goal function if we set $p = B(E)$

(equations 4 and 5 with $W(E) = w(B(E)) = w(p)$). Then, by the uniqueness mentioned, they must be identical. This reasoning shows that Corollary 4 holds for all proper scoring rules and not just for the quadratic one. It confirms that B rather than W is a non-Bayesian analogue of subjective probability.

In practice, we first infer the (for the participant) optimal $R(p)$ for a set of objective probabilities p which is so dense that we obtain a sufficiently accurate estimation of R and R^{-1} . In our experiment, we will consider all values $p = j/20$ for $j \geq 10$. Then, for all uncertain events E (or E^c if $r < 0.5$), we derive $B(E)$ from the observed r_E using equation (12). For $r_E = 0.5$, $B(E)$ and the inverse p may not be uniquely determined because of the flat part of $R_{\text{non-EU}}$ in Figure 1. The case $r < 0.5$ follows from equations (2) and (12), as always. We call the function R^{-1} the *risk correction* (for proper scoring rules). $R^{-1}(r_E) = B(E)$ is the corrected reported probability.

6. AN ILLUSTRATION OF OUR MEASUREMENT OF BELIEF

This section describes risk corrections for a participant in the experiment to illustrate how our method can be applied empirically. It will show that Corollary 4 is the only result of the theoretical analysis needed to apply our method. Results and curves for $r < 0.5$ are derived from $r > 0.5$ using equation (2); we will not mention this point explicitly in what follows.

The left side of Figure 2 displays the performance of stock 12 (the Royal Begemann Group) in our experiment from 1 January, 1991 until 1 June, 1991, as given to the participants. Further details (such as the absence of a unit on the y-axis) will be explained in Section 7. The right side of the figure displays two disjoint intervals S and T , and their union $I = S \cup T$. For each of the intervals S, T and I , participants reported the probability of the stock ending up




FIGURE 2

Layout of the screens (with braces and letters added)

in that interval on 1 January, 1992 (with some other questions in between the three questions considered here). For participant 14, the results are as follows.

$$r_S = 0.35; r_T = 0.55; r_I = 0.65. \quad (13)$$

Under additivity of reported probability, $r_S + r_T - r_I$ (the *additivity bias*, defined in general following equation (21)), should be 0, but here it is not and additivity is violated.

$$\text{The additivity bias is } 0.35 + 0.55 - 0.65 = 0.25. \quad (14)$$

Table 1 and Figure 3 (in inverted form) display the reported probabilities that we measured from this participant as a function of objective probabilities, with the curves explained later. We use progressive averages (midpoints between data points) so as to reduce noise.⁴

For simplicity of presentation, we analyse the data here using linear interpolation. Then $R(0.23) = 0.35$.⁵ Using this value for $R(0.23)$, using the values $R(0.56) = 0.55$, and $R(0.77) = 0.65$, and, finally, using equation (12), we obtain the following corrected reported probabilities:

$$B(S) = R^{-1}(0.35) = 0.23; B(T) = R^{-1}(0.55) = 0.56; B(I) = R^{-1}(0.65) = 0.77; \\ \text{the additivity bias is } 0.23 + 0.56 - 0.77 = 0.02. \quad (15)$$

The risk correction has reduced the violation of additivity, which, according to Bayesian principles, can be interpreted as a desirable move towards rationality. In the experiment described in the following sections, we will see that this effect is statistically significant for single evaluations (treatment “ $t = \text{ONE}$ ”), but is not significant for repeated payments and decisions (treatment “ $t = \text{ALL}$ ”).

It is statistically preferable to fit data with smoother curves than those resulting from linear interpolation. We derived “decision-theoretic” parametric curves for $R(p)$ from Corollary 3, with further assumptions explained in Section 8.1.⁶ The resulting curve for participant 14 is given in Figure 3. The equality $B = R^{-1}(r)$ and this curve imply:

$$B(S) = R^{-1}(0.35) = 0.24; B(T) = R^{-1}(0.55) = 0.59; B(I) = R^{-1}(0.65) = 0.76; \\ \text{the additivity bias is } 0.24 + 0.59 - 0.76 = 0.07. \quad (16)$$

The uncorrected additivity bias is, again, reduced. For this participant, the quadratic curve, explained in Section 12, happens to be indistinguishable from the decision-theoretic curve.

TABLE 1
Progressive average reported probabilities $R(p)$ of participant 14

P	0.025	0.075	0.125	0.175	0.225	0.275	0.325	0.375	0.425	0.475
$R(p)$	0.067	0.192	0.267	0.305	0.345	0.382	0.422	0.435	0.437	0.470
P	0.525	0.575	0.625	0.675	0.725	0.775	0.825	0.875	0.925	0.975
$R(p)$	0.530	0.563	0.565	0.578	0.618	0.655	0.695	0.733	0.808	0.933

4. For each midpoint between two given probabilities p , we calculated the average report for the adjacent probabilities. For instance, to compute the $R(p)$ for $p = 0.625$, we averaged the reported probabilities for $p = 0.6$ and those for $p = 0.65$.

5. We have $0.23 = 0.865 \times 0.225 + 0.135 \times 0.275$, $R(0.225) = 0.345$, and $R(0.275) = 0.382$, so that $R(0.23) = R(0.865 \times 0.225 + 0.135 \times 0.275) = 0.865 \times R(0.225) + 0.135 \times R(0.275) = 0.865 \times 0.345 + 0.135 \times 0.382 = 0.35$.

6. The decision-theoretic curve in the figure is the function $p = B(E) = \frac{r}{r + (1-r) \frac{0.26(1-(1-r)^2) - 1.26}{0.26(1-r^2) - 1.26}}$ on $(0.5, 1)$, in agreement with Corollaries 3 and 4, where we estimated $w(p) = p$ and found $\rho = -0.26$ as the optimal value for $U(x)$ in equation (17).

7. AN EXPERIMENTAL APPLICATION OF RISK CORRECTIONS: METHOD

The following five sections (Sections 7–11) present the third part of this paper. These sections give an experimental implementation of our new measurement method. We first describe the two main treatments in detail. Section 10 presents a third, control, treatment.




FIGURE 3
 $R^{-1}(r)$ of participant 14

7.1. Participants

For the first two treatments, $N = 93$ students from a wide range of disciplines (45 economics, 13 psychology and 35 other disciplines) participated in the experiment. They were self-selected from a mailing list of approximately 1100 people.

7.2. Procedure

Participants were seated in front of personal computers in six groups of approximately 16 participants each. They first received an explanation of the QSR, given in Appendix D. Then, for each uncertain event, participants could first report a probability (in percentages) by typing in an integer from 0 to 100. Subsequently, the confirmation screen displayed a list box with probabilities and the corresponding score when the event was true/not true, illustrated in Figure 4.

All the figures (including Figure 2) are reproduced here in black and white; in the experiment, we used colours to further clarify the figures. The entered probability and the corresponding score were preselected in this list box. The participant could confirm the decision or change to another probability by using the up or down arrow or by scrolling to another probability using the mouse. The event itself was also visible on the confirmation screen. Thus, the reported probability r finally resulted for the uncertain event.

Probability	Your score if statement is true	Your score if statement is not true
27%	4671	9271
28%	4816	9216
29%	4959	9159
30%	5100	9100
31%	5239	9039
32%	5376	8976
33%	5511	8911
34%	5644	8844
35%	5775	8775
36%	5904	8704

send

FIGURE 4
A confirmation screen

7.3. Stimuli

The participants provided 100 reported probabilities r for events with unknown probabilities in the *stock price part* of the experiment. For these events, we fixed 1 June, 1991, as the “evaluation date”. The uncertain events always concerned the question whether or not the price of a stock would lie in a target interval 7 months after the evaluation date. For each stock, the participants received a graph depicting the price of the stock on 0, 1, 2, 3, 4 and 5 months before the evaluation date, as well as an upper and lower bound to the price of the stock on the evaluation date. Figure 2 (without the braces and letters) gives an example of the layout. We used 32 different stocks, all real-world stock market data from the 1991 Amsterdam stock exchange. After four practice questions, the graph of each stock price was displayed once in the questions 5–36, once in the questions 37–68 and once in the questions 69–100. We thus obtained three probabilistic judgements of the performance of each stock, once for a large target interval and twice for the small target intervals that partitioned the large target interval (Figure 2). We partially randomized the order of presentation of the elicitations. Each stock was presented at the same place in the first, second and third 32-tuple of elicitations, so as to ensure that questions pertaining to the same stock were always far apart. The order of presentation of the one large and the two small intervals for each stock was not randomized stochastically, but was varied systematically, so that all orders of big and small intervals occurred equally often. We also maximized the variation of whether small intervals were both very small, or were both moderately small, or one was very small and one was moderately small.

In the *calibration part* of the experiment, participants essentially made the same decisions as in the stock price part, but now for 20 events with objective probabilities. We used two 10-sided dice to determine the outcome of the different prospects. One die determined the first digit and the other determined the second digit of a random number below 100. An event with probability 0.25 was, for example, described as “The outcome of the roll with two 10-sided dice is in the range 01–25”. The subjects then chose r in $(1 - (1 - r)^2)_E(1 - r^2)$ with E the event as described. This amounts to choosing the optimal r for objective probability distributions $(1 - (1 - r)^2)_{0.25}(1 - r^2)$ in the example. We obtained measurements of the reported probabilities corresponding to the objective probabilities 0.05, 0.10, 0.15, . . . , 0.85, 0.90 and 0.95 (we measured the objective probability 0.95 twice). The decision screen was very similar to Figure 4, but we wrote “row-percentage” instead of “probability” and “your score if the roll of the die is 01–25” instead of “your score if the statement is true;” and so on.

7.4. *Motivating participants*

Depending on whether the uncertain event obtained or not, and on the reported probability for the uncertain event, a number of points were determined for each question by means of the QSR (equation 1), using 10,000 points as the unit of payment so as to have integer scores with four digits of precision. Thus, the maximum score for one question was 10,000, the minimum score was 0 and the certain score resulting from reported probability 0.5 was 7500 points.

In treatment $t = \text{ALL}$, the sum of all points for all questions was calculated for each participant and converted to money using an exchange rate of 60,000 points = €1, yielding an average payment of €15.05 per participant. For the calibration part, we then used a box with 20 separate compartments containing pairs of 10-sided dice to determine the outcome of each of the 20 prospects at the same time for the treatment $t = \text{ALL}$.

In treatment $t = \text{ONE}$, the random incentive system was used. That is, at the end of the experiment, 1 out of the 120 questions that they answered was selected at random for each participant, and the points obtained for this question were converted to money using an exchange rate of 500 points = €1, yielding an average payment of €15.30 per participant.

All payments were done privately at the end of the experiment.

7.5. *Analysis*

For the calibration part, we only need to analyse probabilities of 0.5 or higher, using equation (3) (see also Observation A2). Every observation for $p < 0.5$ amounts to an observation for $p' = 1 - p > 0.5$. It implies that we have two observations for all $p > 0.5$ (and three for $p = 0.95$).

We first analyse the data at the group level, assuming homogeneous participants. We start from general probabilistic sophistication. Note that this model can be estimated using a non-parametric procedure. If the agent is willing to go through a large series of correction questions, it is possible to measure the corresponding reported probability of each objective probability repeatedly. In this way, an accurate estimate of the whole correction curve can be obtained without making assumptions about the utility function or the weighting function. This procedure is appropriate if the goal is to correct an expert, for example, correct the reports provided by a weatherman. In applications of experimental economics where subjects participate for a limited amount of time, the researcher will only be able to collect a limited number of observations of the correction curve. Then it is more appropriate to follow a parametric approach to elicit the curve that best fits the observations. In this paper, we used parametric fittings. For U , we used the *power utility with parameter* ρ , also known as the family of constant relative risk aversion (CRRA),⁷ and the most popular parametric family for fitting utility, which is defined as follows:

$$\begin{aligned} \text{For } \rho > 0: U(x) &= x^\rho; \\ \text{For } \rho = 0: U(x) &= \ln(x); \\ \text{For } \rho < 0: U(x) &= -x^\rho. \end{aligned} \tag{17}$$

It is well known that the unit of payment is immaterial for this family. The most general family that we consider for $w(p)$ is Prelec's (1998) two-parameter family

$$w(p) = (\exp(-\beta(-\ln(p))^\alpha)), \tag{18}$$

7. We avoid the latter term because, in non-expected utility models as relevant for this paper, risk aversion depends not only on the curvature of utility.

chosen for its analytic tractability and good empirical performance. We will mostly use the one-parameter subfamily with $\beta = 1$, as in equation (7), for reasons explained later. Substituting the above functions yields

$$B(E) = \exp \left(- \left(\frac{-\ln \left(\frac{r(2r-r^2)^{1-\rho}}{(1-r)(1-r^2)^{1-\rho} + r(2r-r^2)^{1-\rho}} \right)}{\beta} \right)^{1/\alpha} \right)$$

for equation (10).

The model we estimate for each subject separately is as follows:

$$R_k(j/20) = h(j/20, \alpha, \rho) + \varepsilon_k(j/20). \quad (19)$$

Here, $R_k(j/20)$ is the reported probability of the participant for known probability $p = j/20$ ($10 \leq j \leq 19$) in treatment t ($t = \text{ALL}$ or $t = \text{ONE}$) for the k -th measurement for this probability, with only $k = 1$ for $j = 10$, $k = 1, 2$ for $11 \leq k \leq 18$, and $k = 1, 2, 3$ for $j = 19$. With β set equal to 1, α is the remaining probability weighting parameter (equation 18), and ρ is the power of utility (equation 17). The function h is the inverse of equation (11). Although we have no analytic expression for this inverse, we could calculate it numerically in the analyses. The error terms $\varepsilon_k(j/20)$ are drawn from a truncated normal distribution with mean 0 and variance σ^2 . The distribution of the error terms is truncated because reported probabilities below 0 and above 1 are excluded by design. Error terms are identically and independently distributed across choices. We employed maximum likelihood to estimate the parameters of equation (19).

We also carried out an analysis at the aggregate level of the calibration part, with α_t and ρ_t , that is, with these parameters depending on the treatment but not on the participant. To correct for individual differences, we added an individual-specific constant $c_{s,t}$ to the equation where s refers to the participant and t to the treatment:

$$R_{s,t,k}(j/20) = h(j/20, \alpha_t, \rho_t) + c_{s,t} + \varepsilon_{s,t,k}(j/20, \sigma_t^2). \quad (20)$$

Here, the error terms are independent across subjects, treatments and choices.

In the stock price part, violations of additivity were tested. With I the large interval of a stock, being the union $S \cup T$ of the two small intervals S and T , additivity of the uncorrected reported probabilities implies

$$r_S + r_T - r_I = 0. \quad (21)$$

Hence, $r_S + r_T - r_I$ is an index of deviation from additivity, which we call the *additivity bias* of r .

Under the null hypothesis of additivity for corrected reported probabilities B , binary additivity holds, and we can obtain $B(S) = 1 - B(S^c)$ for small intervals S in the experiment (cf. equation 2). Thus, under additivity of B , we have

$$B(S) + B(T) - B(I) = 0. \quad (22)$$

Hence, $B(S) + B(T) - B(I)$ is an index of deviation from additivity of B , and is B 's *additivity bias*.

We next discuss tests of the additivity bias. For each individual stock, and also for the average over all stocks, we tested for both treatments $t = \text{ONE}$ and $t = \text{ALL}$: (i) whether the additivity bias was zero or not, both with and without risk correction; (ii) whether the average additivity bias, as relevant for aggregated group behaviour and expert opinions, was enlarged or reduced by correction; (iii) whether the absolute value of the additivity bias, as relevant for additivity at the individual level, was enlarged or reduced by correction. We report only the tests for averages over all stocks.

8. RESULTS OF THE CALIBRATION PART

Risk corrections and, in general, QSR measurements, do not make sense for participants who are hardly responsive to probabilities, so that $R(p)$ is almost flat on its entire domain. Hence, we kept only those participants for whom the correlation between reported probability and objective probability exceeded 0.35. We thus dropped four participants. The following analyses are based on the remaining 89 participants.

8.1. Group averages

We did several tests using equation (18) with β as a free (treatment dependent or treatment independent) variable, but β 's estimates added little extra explanatory power to the other parameters and usually were close to 1. Hence, we chose to focus on a more parsimonious model in which the restriction $\beta_{\text{ONE}} = \beta_{\text{ALL}} = 1$ is employed. Table 2 lists the estimates for the model of equation (20) for $\beta = 1$ (equation 7 instead of equation 18), together with the estimates of some models with additional restrictions. We first give results at the aggregate level. Because there turns out to be a strong correlation between the α and the ρ parameters, estimation results where both parameters are estimated simultaneously cannot be trusted, and we only report the results where either α or ρ is estimated.

8.1.1. Overall need for risk correction. The first row of Table 2 shows the results without any correction. The second row presents the results when utility curvature is introduced. The likelihood improves significantly (likelihood ratio test, $p = 0.01$) and substantially, so that risk correction is called for. Risk correction can also be done by probability weighting. This is done in the third row of the table. Probability weighting also increases the likelihood of observing the data significantly compared with the model without correction, but less so than utility curvature does. Therefore, in the remainder of the paper, we focus on risk correction obtained through utility curvature.

8.1.2. Comparing the two treatments. At the aggregate level, risk correction is needed less in treatment ALL than in treatment ONE, as the fourth and fifth rows show. In treatment

TABLE 2
Estimation results at the aggregate level

Row	Restrictions	σ_{ONE}	α_{ONE}	ρ_{ONE}	σ_{ALL}	α_{ALL}	ρ_{ALL}	$-\log L$
1	$\alpha_{\text{ONE}} = \alpha_{\text{ALL}} =$ $\rho_{\text{ONE}} = \rho_{\text{ALL}} = 1$	9.00** (0.21)	—	—	8.36** (0.20)	—	—	6373.21
2	$\alpha_{\text{ONE}} = \alpha_{\text{ALL}} = 1$	8.73** (0.20)	—	0.43** (0.09)	8.36** (0.20)	—	0.94** (0.07)	6345.43
3	$\rho_{\text{ONE}} = \rho_{\text{ALL}} = 1$	8.82** (0.21)	0.69** (0.03)	—	8.35** (0.20)	1.09** (0.07)	—	6354.14
4	$\alpha_{\text{ONE}} = \alpha_{\text{ALL}} =$ $\rho_{\text{ONE}} = 1$	9.00** (0.21)	—	—	8.36** (0.20)	—	0.94** (0.07)	6372.87
5	$\alpha_{\text{ONE}} = \alpha_{\text{ALL}} =$ $\rho_{\text{ALL}} = 1$	8.73** (0.20)	—	0.43** (0.09)	8.36** (0.20)	—	—	6345.77
6	$\alpha_{\text{ONE}} = \alpha_{\text{ALL}} =$ $1, \rho_{\text{ONE}} = \rho_{\text{ALL}}$	8.78** (0.21)	—	0.70** (0.06)	8.41** (0.20)	—	—	6556.48

Notes: Standard errors in parentheses, ** denotes significance at the 1% level.




FIGURE 5

Corrected versus uncorrected probability

ONE, the likelihood is improved significantly (compare the fifth and the first row; likelihood ratio test, $p = 0.01$) but in treatment ALL the likelihood is not improved significantly (compare the fourth and the first row; likelihood ratio test, $p > 0.10$). We obtain $\rho_{\text{ONE}} < \rho_{\text{ALL}}$: if only one decision is paid out, then participants exhibit more concave curvature of utility than when all decisions are paid out. Given the same degree of probability weighting, it implies more risk aversion for $t = \text{ONE}$ than for $t = \text{ALL}$ (and R closer to 0.5). The finding is supported by comparing the sixth row of Table 2 with the restriction $\rho_{\text{ONE}} = \rho_{\text{ALL}}$, to the second row. This restriction significantly reduces the likelihood of observing the data (likelihood ratio test, $p = 0.01$).

Figure 5, on the basis of the estimates reported in the second row of Table 2, displays the resulting average risk correction for the two treatments separately. The figure illustrates that risk correction is clearly needed at the aggregate level in treatment ONE.

8.2. Individual analyses

8.2.1. Need for risk correction at the individual level. There is considerable heterogeneity in each treatment. Whereas the corrections required were small at the level of group averages, they are big at the individual level. This appears from Figure 6, which displays the cumulative distribution of the (per-subject) estimated ρ -coefficients for each treatment, assuming $\alpha = \beta = 1$. (The figure also displays a treatment $t = \text{ALLnp}$ that will be explained in Section 10.) There are wide deviations from the value $\rho = 1$ (i.e. no correction) on both sides. As seen from the group-average analysis, there are more deviations at the risk-averse side of $\rho < 1$.

8.2.2. Comparing the two treatments. The ρ -coefficient distribution of treatment $t = \text{ONE}$ dominates the ρ -coefficient distribution of treatment $t = \text{ALL}$. Thus, the ρ -coefficients for $t = \text{ONE}$ are lower than for $t = \text{ALL}$ ($p = 0.001$: two-sided Mann–Whitney test). It confirms the result from Table 2 that there is more risk aversion for group averages, moving R in the direction of 0.5, for $t = \text{ONE}$ than for $t = \text{ALL}$. The figure also shows that, in an absolute sense, there is more deviation from $\rho = 1$ for $t = \text{ONE}$ than for $t = \text{ALL}$, implying that there

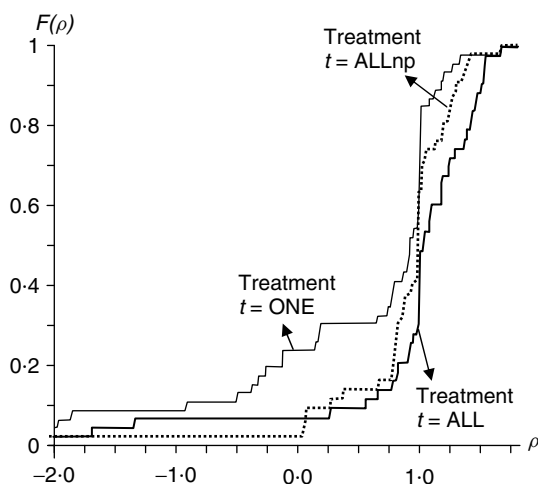


FIGURE 6
Cumulative density of ρ

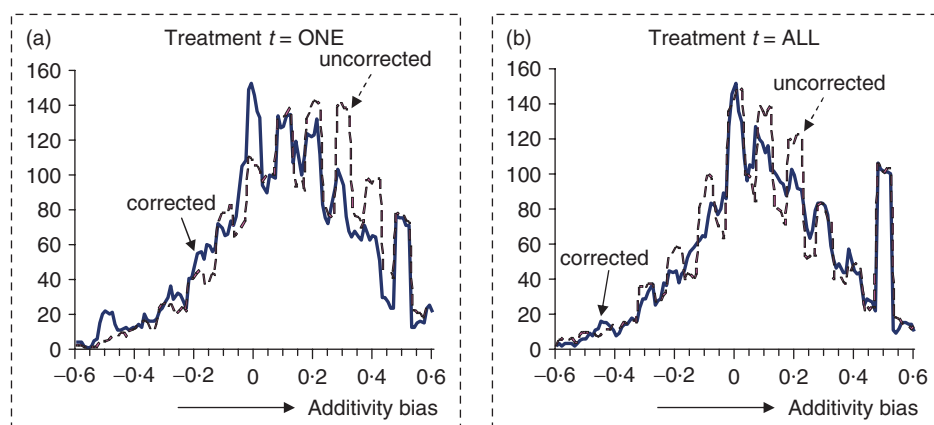
are more deviations from expected value and more risk corrections for $t = \text{ONE}$ than for $t = \text{ALL}$.

Unlike the median ρ -coefficients that are fairly close to each other for the two treatments (0.92 for $t = \text{ONE}$ versus 1.04 for $t = \text{ALL}$), the mean ρ coefficients are substantially different (0.24 for $t = \text{ONE}$ versus 0.91 for $t = \text{ALL}$), which is caused by skewedness to the left for $t = \text{ONE}$. That is, there is a relatively high number of strongly risk-averse participants for $t = \text{ONE}$. Analyses of the individual ρ parameters (two-sided Wilcoxon signed rank sum tests) confirm findings of group-average analyses in the sense that the ρ -coefficients are significantly smaller than 1 for $t = \text{ONE}$ ($z = -3.50$, $p = 0.0005$), but not for $t = \text{ALL}$ ($z = 1.42$, $p = 0.16$).

9. RESULTS FOR THE STOCK PRICE PART: RISK CORRECTION AND ADDITIVITY

All comparisons in this section are based on two-sided Wilcoxon signed rank sum tests. Figure 7 displays the data, aggregated over both stocks and individuals, of the additivity biases for $t = \text{ONE}$ and for $t = \text{ALL}$. The figures show that the additivity bias is more often positive than negative, in agreement with common findings in the literature (Tversky and Koehler, 1994). Indeed, for virtually all stocks the additivity bias is significantly positive for both treatments, showing in particular that additivity does not hold. This also holds when taking the average additivity bias over all stocks as one data point per participant ($z = 5.27$, $p < 0.001$ for $t = \text{ONE}$, $z = 4.35$, $p < 0.001$ for $t = \text{ALL}$). We next consider whether risk corrections reduce the violations of additivity.

Let us first consider treatment $t = \text{ONE}$. Here, the risk corrections reduce the average additivity bias significantly for 27 of the 32 stocks and enlarge it for none. We only report the statistics for the average additivity bias over all stocks taken as one data point per participant, which has overall averages 0.163 (uncorrected) and 0.120 (corrected), with the latter significantly smaller ($z = 3.21$, $p = 0.001$). For assessing the degree of irrationality (additivity violation) at the individual level, the absolute values of the additivity bias are relevant. For $t = \text{ONE}$, Figure 7 suggests that these are smaller after correction, because on average the



For each interval $[\frac{j-2.5}{100}, \frac{j+2.5}{100}]$ of length 0.05 around $\frac{j}{100}$, we counted the number of additivity biases in the interval, aggregated over 32 stocks and 89 individuals, for both treatments. With risk-correction, there were 65 additivity biases between 0.375 and 0.425 in the treatment $t = ONE$, and without risk correction there were 95 such; and so on.

FIGURE 7

Empirical density of additivity bias for the two treatments

corrected curve is closer to 0 on the x -axis. These absolute values were significantly reduced for nine stocks, and enlarged for none. Again, we only report the statistics for the average absolute value of the additivity bias over all stocks taken as one data point per participant, which has overall averages 0.239 (uncorrected) and 0.228 (corrected), with the latter significantly smaller ($z = 2.26$, $p = 0.02$).

For $t = ALL$, risk corrections did not significantly alter the average additivity bias. More specifically, it gave a significant increase for three stocks and a significant decrease for one stock, which, for 32 stocks, suggests no systematic effect. The latter was confirmed when we took the average additivity bias over all stocks for each individual, with no significant differences generated by correction (average 0.128 uncorrected and average 0.136 corrected; $z = -1.64$, $p = 0.1$). Similar results hold for absolute values of additivity biases, which gave a significant increase for one stock and a significant decrease for no stocks. Taking the average additivity bias over all stocks as one data point per participant (average 0.237 uncorrected and average 0.239 corrected; $z = -0.36$, $p = 0.70$) also gave no significant difference.

Risk correction reduces the additivity bias for treatment $t = ONE$ to a level similar to that observed for $t = ALL$ (averages 0.120 and 0.136). The overall pattern is that beliefs for $t = ONE$ after correction, and for $t = ALL$ both before and after correction, exhibit a similar degree of violation of additivity, which is clearly different from zero. The additivity bias is not completely caused by non-linear risk attitudes when participants report probabilities, but has a genuine basis in beliefs.

10. A TREATMENT WITHOUT EXPLICIT REFERENCE TO BELIEFS OR PROBABILITY

This section briefly reports the results of a robustness check of our experimental design. In agreement with the current practice of scoring rules, the instructions of our main treatments repeatedly used the terms probability and belief. These terms may have influenced the subjects.

To assess such influences, we performed a control treatment in which we did not refer to probabilities or beliefs.⁸ Thus, in Figure 4 we now used the expression “choose a number” instead of probability. In the instructions of this control treatment, we similarly asked subjects to choose numbers without calling them probabilities, and we dropped all interpretations of likelihood. In this manner, we ran the control treatment for $t = \text{ALL}$. We chose $t = \text{ALL}$ rather than $t = \text{ONE}$ because the former is most commonly used in applications of proper scoring rules. We refer to the control treatment as $t = \text{ALLnp}$ (np for no probabilities). $N = 44$ students participated. The number of participants dropped from the analysis because their correlation between reported and objective probability was below 0.35, was now 2. In all other respects, the new treatment was identical to the $t = \text{ALL}$ treatment in the main experiment.

The results confirmed all patterns and inequalities found for $t = \text{ALL}$. We give some numerical details for individual analyses. The ρ 's of $t = \text{ALLnp}$ are not significantly different from those of $t = \text{ALL}$ ($z = 1.57$, $p = 0.12$), with a similar median (1.00 for $t = \text{ALLnp}$ vs. 1.04 for $t = \text{ALL}$) and mean (0.80 for $t = \text{ALLnp}$ vs. 0.91 for $t = \text{ALL}$). They, accordingly, are not significantly below 1 either ($z = 0.52$, $p = 0.60$), and they also exceed the ρ for $t = \text{ONE}$ ($z = -2.30$, $p = 0.02$).

The additivity bias is, again, positive, showing that additivity is violated, for most individual stocks. It also is when taking the average additivity bias over all stocks per participant ($z = 4.47$, $p < 0.001$). Risk corrections did not significantly increase or decrease the average additivity bias for any stock. For the (absolute) average additivity bias over all stocks per participant, we again found no significant difference between the non-corrected and corrected average additivity bias ($z = 0.378$, $p = 0.71$; $z = 0.265$, $p = 0.79$ for absolute values). The risk-corrected average additivity bias for $t = \text{ALLnp}$ being virtually the same as for $t = \text{ALL}$ (0.126 vs. 0.136) obviously implies that it is also equal to the one for $t = \text{ONE}$ (0.120). In summary, all the results for the $t = \text{ALL}$ treatment were confirmed by the $t = \text{ALLnp}$ treatment, suggesting that the explicit use of the term probability in our instructions did not alter the results.

11. DISCUSSION OF EXPERIMENT

11.1. *Methods*

We chose the evaluation date (1 June, 1991) sufficiently long ago to ensure that participants would be unlikely to recognize the stocks or have private information about them. In addition, no numbers were displayed on the vertical axis, making it extra hard for participants to recognize specific stocks. We thus ensured that participants based their probability judgements entirely on the prior information about past performance of the stocks given by us. Given the large number of questions, it is unlikely that participants noticed that the graphs were presented more than once (three times) for each stock. Indeed, in informal discussions after the experiment no participant showed awareness of this point.

In some studies in the literature, the properness of scoring rules is explained to participants by stating that it is in their best interest to state their true beliefs, either without further explanation, or with the claim added that they will thus maximize their “expected” money. A drawback of this explanation is that expected value maximization is empirically violated, which is the central topic of this paper (Section 3), so that the recommendation is debatable. We, therefore, used an alternative explanation that relates properness of one-off events to observed frequencies of repeated events (Appendix D).

8. This treatment was recommended to us by a referee and the editor.

11.2. *Optimal incentive scheme*

After some theoretical debates about the random incentive system (Holt, 1986), as in our treatment $t = \text{ONE}$, the system was tested empirically and found to be incentive compatible (Starmer and Sugden, 1991; Lee, 2008). It is today the almost exclusively used incentive system for the measurement of individual preferences (Myagkov and Plott, 1997; Harrison, Lau and Williams, 2002; Holt and Laury, 2002). Unlike repeated payments, it avoids income effects such as Thaler and Johnson's (1990) house money effect, and the drift towards expected value and linear utility that is commonly generated by repeated choice.⁹ For the purpose of measuring individual preference, the treatment $t = \text{ONE}$ is, therefore, preferable. When the purpose is, however, to derive subjective probabilities from proper scoring rules, and no risk correction is possible, then a drift towards expected value is actually an advantage, because uncorrected proper scoring rules assume expected value. This point agrees with our findings, where less risk correction was required for the $t = \text{ALL}$ treatment. Li (2007) discussed other arguments for and against repeated rewarding when events are not verifiable and binary rewards have to be used.

For some applications, group averages of probability estimates are most relevant, such as when aggregating expert judgements or predicting group behaviour. Then our statistical results regarding "non-absolute" values of reported probabilities are most relevant. For the assessment of rationality at the individual level, absolute values of the additivity biases are most relevant.

11.3. *Choice of parameters*

The lack of extra explanatory power of parameter β in equation (18) should come as no surprise because β and α imply similar phenomena on $[0.5, 1]$, enhancing risk aversion there. They mainly deviate from one another on $[0, 0.5]$, where β continues to enhance risk aversion but α enhances the inverse-S shape that is mostly found empirically. The domain $[0, 0.5]$ is, however, not relevant to our study (Observation A2).

11.4. *Pragmatic applications*

More tractable families can be used to fit the reported probabilities than the decision-theory-based curves that we used. For example, in Figure 3, we also used quadratic regression to find the curve $p = a + br + cr^2$ that best fits the data. For most participants, the curve is virtually indistinguishable from the decision-theoretic curve. This observation, together with Corollary 4 which demonstrates that we only need the readily observable reported probabilities and not the actual utility function or probability weighting function to apply our method, shows that applications of our method are straightforward. The theoretical analysis of this paper, and the decision-theory-based curve fitting that we adopted, served to prove that our method is in agreement with modern decision theories. If this thesis is accepted, and the only goal is to obtain corrected reported probabilities, then one may choose the pragmatic shortcuts just described.

11.5. *General discussion*

We emphasize that the biases because of violations of expected value that we correct for need not concern mistakes or irrationalities in decision making. Deviations from risk neutrality need not be irrational and, according to some, even deviations from Bayesian beliefs need not be

9. It is required that the repeated choices are perceived as sufficiently uncorrelated. Correlation can enhance the perception of, and aversion to, ambiguity (Halevy and Feltkamp, 2005).

irrational, nor need the corresponding ambiguity attitudes be irrational (Gilboa and Schmeidler, 1989; Schmeidler, 1989). The required corrections concern empirical deficiencies of the model of expected value, that is, they concern biases on the part of the researchers analysing the data.

Under proper scoring rules, beliefs are derived solely from decisions, and equation (1) is taken purely as a decision problem, where the only goal of the agent is to optimize the prospect received. We considered two treatments that explicitly referred to probabilities and beliefs, and a treatment that did not do so, finding no differences between the two. Thus, this paper has analysed proper scoring rules purely from the decision-theoretic perspective supported with real incentives and has corrected only for biases resulting therefrom. Many studies have investigated direct judgements of belief without real incentives, and then many other aspects play a role, leading for instance to the often found overconfidence. Such introspective effects are beyond the scope of this paper.

An immediate advantage of our calibration measurement, prior to any theoretical analysis, is that it helps to identify subjects whose understanding of the concepts to be measured is below what is minimally acceptable. Indeed, subjects for whom the correlation between objective and reported probabilities is very low clearly have little clue what likelihood means. Their reported probabilities are of so little interest that we recommend dropping them from the sample. If we are interested in the beliefs of such subjects in more elaborate studies, then further teaching and learning will be called for.

The experimental data show that a substantial correction of reported probabilities needs to be made for a subset of the subjects. The fraction of the sample that needs substantial corrections is larger when only a single large-stake decision is paid than when repeated small decisions are paid. Our conclusion is that it is desirable to correct agents' reported probabilities elicited with scoring rules, especially if only a single large-stake decision is paid. A drawback of our method is, obviously, that it requires individual measurements of QSRs for given probabilities. If it is not possible to obtain individual measurements of the correction curve, then it will be useful to use best guess corrections: for instance, through averages obtained from individuals as similar as possible. Thus, at least, the systematic error for the group average to risk attitude has been corrected for as well as is possible without requiring extra measurements. In this respect, the average curves in Figure 5 are reassuring for existing studies, because these curves suggest that only small corrections were needed for the group averages in our context.

Several methods have been used in the literature to measure the subjective degree of belief of an agent in an event E . Mostly these have been derived from: (i) binary preferences, which only give inequalities or approximations; (ii) binary indifferences, which are hard to elicit, for example, by the complex Becker–DeGroot–Marschak mechanism (Karni and Safra, 1987; Braga and Starmer, 2005) or bisection (Abdellaoui, Vossman and Weber, 2005); (iii) introspection, which is not revealed preference based let alone incentive compatible. Proper scoring rules provide an efficient way to measure subjective beliefs while avoiding the problems mentioned.

12. THEORETICAL DISCUSSION

A way to reveal $B(E)$ from observed choice, as an alternative to our method, is by revealing the *matching probability* p of event E , defined through the equivalence

$$x_p y \sim x_E y \quad (23)$$

for some preset $x > y$, say $x = 100$ and $y = 0$. Then $w(B(E))(U(x) - U(y)) = w(p)(U(x) - U(y))$, and $B(E) = p$ follows. Wakker (2004) discussed the interpretation of equations (9) and

(23) as belief. Matching probabilities were commonly used in early decision analysis (Raiffa, 1968, section 5.3; Yates, 1990, pp. 25–27) under the assumption of expected utility. A recent experimental measurement is in Holt (2006, chapter 27), who also assumed expected utility. Abdellaoui, Vossman and Weber (2005) measured and analysed them in terms of prospect theory, as does our paper. A practical difficulty is that the measurement of matching probabilities requires the measurement of indifference, and these are not easily inferred from choice. For example, Holt (2006) used the Becker–DeGroot–Marschak mechanism, and Abdellaoui, Vossman and Weber (2005) used a bisection method. If we want to measure $B(E)$ for only few events E , then matching probabilities provide a tractable alternative to our correction method. Our method is more efficient when dealing with many events. The measurement of the correction curve is a one-time investment that can next be applied to an unlimited number of events.

Another way to correct reported probabilities is through calibration. Then many reported probabilities are collected over time and are related to observed relative frequencies. Calibration has been studied in game theory (Sandroni, Smorodinsky and Vohra, 2003) and has been applied to weather forecasters (Murphy and Winkler, 1974). It needs extensive data, which is especially difficult to obtain for rare events such as earthquakes, and further assumptions such as the stability of the characteristics of these events over time. Clemen and Lichtendahl (2005) discussed these drawbacks and proposed correction techniques for probability estimates in the spirit of our paper, but still based these on traditional calibration techniques. Our correction (“calibration”) technique is considerably more efficient than the traditional ones. It shares with Prelec’s (2004) method the advantage that we need not wait until the truth or untruth of uncertain events has been revealed to implement it.

Allen (1987) proposed to avoid biases of the QSR because of non-linear utility by paying in terms of the probability of winning a prize instead of in terms of money, and this procedure was implemented by McKelvey and Page (1990). The procedure, however, only works if expected utility holds, and there is much evidence against this assumption. Indeed, Selten, Sadrieh and Abbink (1999) showed empirically that payment in probability generates more deviations from risk neutrality than payment in money does.

The decision-based distortion in the direction of 0.5 because of risk aversion in Section 4 is opposite to the overconfidence (probability judgements too far from 0.5) mostly found in direct judgements of probability without real incentives (McClelland and Bolger, 1994), and found among experts seeking to distinguish themselves (Keren, 1991, pp. 224, 252; the “expert bias”, Clemen and Rolle, 2001). Similar optimistic and pessimistic distortions of probability can result from non-linear utility if the probability considered is a consensus probability for a group of individuals with heterogeneous beliefs (Jouini and Napp, 2007).

The curve for non-EU in Figure 1 is flat around $p = 0.5$, or more precisely, on the probability interval $[0.43, 0.57]$. For probabilities from this interval, the risk aversion generated by non-expected utility is so strong that the agent goes for maximal safety and chooses $r = 0.5$, corresponding with the sure outcome 0.75 (cf. Segal and Spivak, 1990; Manski, 2004, footnote 10). Such a degree of risk aversion is not possible under expected utility, where $r = 0.5$ can happen only for $p = 0.5$ (Observation 2). This flat curve cautions against assigning specific levels of belief to observations $r = 0.5$, because proper scoring rules may be insensitive to small changes in the neighbourhood of $p = 0.5$. It in fact means that there the scoring rules, traditionally called proper, are not really proper.

B captures the component of decision making separate from risk attitude. It is common in decision theory to interpret factors separate from risk attitude as ambiguity. Then B reflects ambiguity attitude. There is no consensus about the extent to which ambiguity reflects non-Bayesian beliefs, and to what extent it reflects non-Bayesian decision attitudes separate from

belief. If the equality $B(E) + B(E^c) = 1$ (*binary additivity*) is violated, then it can further be debated whether $B(E)$ or $1 - B(E^c)$ is to be taken as an index of belief (or of ambiguity). Such interpretations have not yet been settled, and further studies are called for. We have mostly referred to B as reflecting beliefs, in order to stay as close as possible to the terminology used today in the literature on proper scoring rules. Irrespective of the interpretation of B , it is clear that the behavioural component of risk attitude should be filtered out before an interpretation of belief can be considered. This paper shows how this filtering out can be done. In Schmeidler (1989), the main paper initiating equations (4) and (5), w was assumed to be linear, with expected utility for given probabilities, and W coincided with B . Schmeidler interpreted this component as reflecting beliefs.

As is common in the mechanism design literature, our correction procedure assumed deterministic choice. A fundamental question concerns how the mechanism performs when agents make mistakes, as in the random utility model (Luce, 1959; McFadden, 1974, 1976). Such mistakes will affect the optimal elicitation procedure. These issues are relevant to the entire mechanism design literature and are a topic for future research.

13. CONCLUSION

This paper has applied modern theories of risk and ambiguity to proper scoring rules. Mutual benefits have resulted for users of proper scoring rules and for analysts of risk and ambiguity. For the former, we have shown which distortions affect their common measurements and how large these distortions are, using theories that are descriptively better than the expected value hypothesis still common today in applications of proper scoring rules. We have provided a procedure to correct for the aforementioned distortions, and a theoretical foundation has been given for interpretations of the resulting measurements as (possibly non-Bayesian) beliefs and ambiguity attitudes. For analyses of risk and ambiguity, we have shown how the remarkable efficiency of proper scoring rules can be used to measure and analyse subjective beliefs and ambiguity attitudes in ways more tractable than is possible with the binary preferences traditionally used.

The feasibility and tractability of our method have been demonstrated in an experiment, where we used it to investigate some properties of beliefs and quadratic proper scoring rules. We found, for instance, that our correction method reduces the violations of additivity in subjective beliefs but does not eliminate them. It confirms that beliefs are genuinely non-Bayesian and that ambiguity attitudes play a central role in proper scoring rules.

APPENDIX A. TECHNICAL REMARKS

For QSR prospects in equation (1), every choice $r < 0$ is inferior to $r = 0$, and $r > 1$ is inferior to $r = 1$. The optimization problem does not change if we allow all real r , instead of $0 \leq r \leq 1$. Hence, solutions $r = 0$ or $r = 1$ can be treated as interior solutions, and they satisfy the first-order optimality conditions.

In general, it may not be possible to derive both w and U from $R(p)$ without further assumptions, that is, U and w may be non-identifiable for proper scoring rules. Under regular assumptions about U and w , however, they have some different implications. The main difference is that, if we assume that U is differentiable (as done throughout this paper) and concave, then a flat part of $R(p)$ around 0.5 must be caused by w (Observation 2).

We next discuss dualities between $B(E)$ and $1 - B(E^c)$ in more detail. Event A is (*revealed*) *more likely* than event B if, for some positive outcome x , say $x = 100$, the agent prefers $x_A 0$ to $x_B 0$. This observation is independent of the outcome $x > 0$. In view of the symmetry of QSRs in Observation 1, for $r \neq 0.5$ the agent will always allocate the highest payment to the most likely of E and E^c . It leads to the following restriction of QSRs.

Observation A1. Under the QSR in equation (1), the highest outcome is always associated with the most likely event of E and E^c .

Hence, QSRs do not give observations about most likely events when endowed with the worst outcome. Similar restrictions apply to all other proper scoring rules considered in the literature so far. It implies the following result.

Observation A2. For the QSR, only the restriction of w to $[0.5, 1]$ plays a role, and w 's behaviour on $[0, 0.5)$ is irrelevant.

For our risk corrections, we need w only on $[0.5, 1]$. An advantage is that the empirical findings about w are uncontroversial on this domain, the general finding being that w underweights probabilities there. This holds both for the mostly found inverse-S shape (Tversky and Kahneman, 1992; Gonzalez and Wu, 1999; Abdellaoui, 2000; Bleichrodt and Pinto, 2000) and for the also often found convex shapes (Goeree, Holt and Palfrey, 2002; van de Kuilen and Wakker 2009).¹⁰

Some details on weak inequalities and corner solutions are as follows. A choice of $r = 0.5$ may be driven by risk aversion, so that no likelihood ordering between E and E^c can then be concluded. A choice of $r \neq 0.5$ (if close to 0.5) may be driven by risk seeking with equal likelihood of E and E^c . Only interior solutions with a strict inequality $r > 0.5$ exclude that E be strictly less likely than E^c .

As with the weighting function w under risk, B is also applied only to the most likely one of E and E^c in the preceding equations, reflecting again the restriction of the QSR of Observation A1. Hence, under traditional QSR measurements, we cannot test binary additivity directly because we measure $B(E)$ only when E is more likely than E^c . These problems can easily be amended by modifications of the QSR. For instance, we can consider prospects

$$(2 - (1 - r)^2)_E(1 - r^2), \quad (\text{A1})$$

that is, QSR prospects as in equation (1) but with a unit payment added under event E . The classical proper scoring rule properties of Section 2 are not affected by this modification, and the results of Section 3 are easily adapted. With this modification, we have the liberty to combine event E with the highest outcome both if E is more likely than E^c and if E is less likely, and we avoid the restriction of Observation A1. We can then observe w of the preceding subsection, and $W(E)$ and $B(E)$ over their entire domain. Similarly, with prospects

$$(1 - (1 - r)^2)_E(2 - r^2), \quad (\text{A2})$$

we can measure the duals $1 - W(E^c)$, $1 - w(1 - p)$, and $1 - B(E^c)$ over their entire domain. In this study, we confined our attention to the QSRs of equation (1) as they are classically applied throughout the literature. We have revealed their biases according to the current state of the art of decision theory, have suggested remedies whenever possible and have signalled the problems that remain. Further investigations of the, we think promising, modifications of QSRs in the preceding equations are left to future studies.

The restrictions of the classical QSRs also hold for the experiment in this paper. There an application of the QSR to events E less likely than their complements are to be interpreted formally as the measurement of $1 - B(E^c)$. The restrictions also explain why the theorems concerned only the case of $r > 0.5$ (with $r = 0.5$ as a boundary solution).

APPENDIX B. PROOFS

For QSR prospects in equation (1), every choice $r < 0$ is inferior to $r = 0$, and $r > 1$ is inferior to $r = 1$. The optimization problem does not change if we allow all real r , instead of $0 \leq r \leq 1$. Hence, solutions $r = 0$ or $r = 1$ can be treated as interior solutions, and they satisfy the first-order optimality conditions.

Proof of Theorem 1. We write π for the decision weight $W(E)$, and consider the general prospects $(a - b(1 - r)^2)_E(a - br^2)$ for any $b > 0$ and $a \in \mathbb{R}$. Theorem 1 concerns the special case of $a = b = 1$. For optimality of interior solutions r , the first-order optimality condition for equation (4) is that

$$\pi U'(a - b(1 - r)^2)2b(1 - r) - (1 - \pi)U'(a - br^2)2br = 0,$$

implying

$$\pi(1 - r)U'(a - b(1 - r)^2) = (1 - \pi)rU'(a - br^2) \quad (\text{B1})$$

10. On $[0, 0.5)$ the patterns is less clear, with both underweighting and overweighting (Gonzalez and Wu, 1999; Abdellaoui, 2000; Bleichrodt and Pinto, 2000).

or

$$\pi U'(a - b(1 - r)^2) = r \times (\pi U'(a - b(1 - r)^2) + (1 - \pi)U'(a - br^2)),$$

and equation (6) follows. \parallel

Proof of Observation 2. If $r = 0.5$ then the marginal utility ratio in equation (6) is 1, and $p = 0.5$ follows. For the reversed implication, assume risk aversion. Then $r > 0.5$ is not possible for $p = 0.5$ because then the marginal utility ratio in equation (6) would be at least 1 so that the right-hand side of equation (6) (where now $W(E) = 0.5$) would at most be 0.5, contradicting $r > 0.5$. Applying this finding to E^c and using equation (2), $r < 0.5$ is not possible either, and $r = 0.5$ follows.

Under strong risk seeking, r may differ from 0.5 for $p = 0.5$. For example, if $U(x) = e^{2.5x}$, then $r = 0.14$ and $r = 0.86$ are optimal, and $r = 0.5$ is a local infimum, as calculations can show. As an aside, the same optimal values of r result under non-expected utility with linear U , and with $w(0.5) = 0.86$. Such large w -values also generate risk seeking. \parallel

Proof of Corollary 2. Let $r > 0.5$ be optimal, and write $\pi = W(E)$. Then equation (B1) implies

$$\pi \times ((1 - r)U'(a - b(1 - r)^2) + rU'(a - br^2)) = rU'(a - br^2),$$

implying

$$\pi = \frac{r}{r + (1 - r) \frac{U'(a - b(1 - r)^2)}{U'(a - br^2)}}. \quad (\text{B2})$$

Applying w^{-1} to both sides yields the corollary. \parallel

In measurements of belief, one first observes r and then derives $B(E)$ from it. Corollary 2 gave an explicit expression. In general, it does not seem to be possible to write r as an explicit expression of $B(E)$ or, in the case of objective probabilities with $B(E) = p$, of the probability p .

Proof of Corollary 4. Theorem 1 implies that the right-hand side of equation (6) is r both as it is, and with p substituted for $B(E)$ (in $W(E) = w(B(E))$). Because equation (6) is strictly increasing in $w(B(E))$, and w is strictly increasing too, $p = B(E)$ follows. \parallel

APPENDIX C. MODELS FOR DECISION UNDER RISK AND UNCERTAINTY

For binary (two-outcome) prospects with both outcomes non-negative, as considered in QSRs, equations (4) and (5) have appeared many times in the literature. References include the early Allais (1953, equation 19.1) and Edwards (1954, figure 3) for risk and, more recently, Luce (1991) for uncertainty. The convenient feature that binary prospects suffice to identify utility U and the non-additive $w(B) = W$ was pointed out by Ghirardato and Marinacci (2001), Gonzalez and Wu (2003), Luce (1991, 2000), Miyamoto (1988), Pfanzagl (1959, p. 287) and Wakker and Deneffe (1996, pp. 1143, 1144–1145).

The convenient feature that most decision theories agree on the evaluation of binary prospects was pointed out by Miyamoto (1988), calling equations (4) and (5) generic utility, and Luce (1991), calling these equations binary rank-dependent utility. It was most clearly analysed by Ghirardato and Marinacci (2001), who called the equations the biseparable model. These three works also axiomatized the model. The agreement for binary prospects was also central in many works by Luce (e.g. Luce, 2000, chapter 3) and in Gonzalez and Wu (2003). Only for more than two outcomes, do the theories diverge (Mosteller and Nogee, 1951, p. 398; Luce, 2000, Introductions to chapters 3 and 5). Theories that also deviate for two outcomes include betweenness models (Chew and Tan, 2005), the variational model (Maccheroni, Marinacci and Rustichini, 2006) and models with underlying multistage decompositions (Halevy and Feltkamp, 2005; Klibanoff, Marinacci and Mukerji, 2005; Nau, 2006; Halevy and Ozdenoren, 2007; Olszewski, 2007).

We next describe some of the agreeing decision theories. Because we consider only non-negative outcomes, losses play no role, and we describe prospect theory only for gains.

We begin with decision under risk, with known objective probabilities $P(E)$. Expected utility (von Neumann and Morgenstern, 1944) is the special case where w is the identity and $B(E) = P(E)$. Kahneman and Tversky's (1979) original prospect theory, Quiggin's (1982) rank-dependent utility and Tversky and Kahneman's (1992) new prospect theory concern the special case of $B(E) = P(E)$, where w now can be non-linear. The case $B(E) = P(E)$ also includes Gul's (1991) disappointment aversion theory.

We next consider the more general case where no objective probabilities need to be given for all events E . Expected utility is the special case where B is an additive, now “subjective”, probability and w is the identity. Choquet expected utility (Schmeidler, 1989) and cumulative prospect theory (Tversky and Kahneman, 1992) start from the general weighting function W , from which B obviously results as $w^{-1}(W)$, with w the probability weighting function for risk. The multiple priors model (Gilboa and Schmeidler, 1989) results with $W(E)$ the infimum value $P(E)$ over all priors P . Under the alpha maxmin model, W is alpha times that infimum plus $1 - \alpha$ times the supremum. Under Machina and Schmeidler’s (1992) probabilistic sophistication, B is an additive probability measure.

APPENDIX D. EXPERIMENTAL INSTRUCTIONS

This appendix is not reproduced here, but is made available online.

Acknowledgements. Glenn Harrison and two anonymous referees made helpful comments.

REFERENCES

- ABDELLAOUI, M. (2000), “Parameter-Free Elicitation of Utilities and Probability Weighting Functions”, *Management Science*, **46**, 1497–1512.
- ABDELLAOUI, M., VOSSMAN, F. and WEBER, M. (2005), “Choice-Based Elicitation and Decomposition of Decision Weights for Gains and Losses Under Uncertainty”, *Management Science*, **51**, 1384–1399.
- ALLAIS, M. (1953), “Le comportement de l’homme rationnel devant le risque: Critique des postulats et axiomes de l’école Américaine”, *Econometrica*, **21**, 503–546.
- ALLEN, F. (1987), “Discovering Personal Probabilities When Utility Functions are Unknown”, *Management Science*, **33**, 542–544.
- ARAGONES, E., GILBOA, I., POSTLEWAITE, A. and SCHMEIDLER, D. (2005), “Fact-Free Learning”, *American Economic Review*, **95**, 1355–1368.
- BERNOULLI, D. (1738), “Specimen Theoriae Novae de Mensura Sortis”, *Commentarii Academiae Scientiarum Imperialis Petropolitanae*, **5**, 175–192.
- BLEICHRODT, H. and PINTO, J. L. (2000), “A Parameter-Free Elicitation of the Probability Weighting Function in Medical Decision Analysis”, *Management Science*, **46**, 1485–1496.
- BRAGA, J. and STARMER, C. (2005), “Preference Anomalies, Preference Elicitation, and the Discovered Preference Hypothesis”, *Environmental and Resource Economics*, **32**, 55–89.
- BRIER, G. W. (1950), “Verification of Forecasts Expressed in Terms of Probability”, *Monthly Weather Review*, **78**, 1–3.
- CAMERER, C. F. and WEBER, M. (1992), “Recent Developments in Modelling Preferences: Uncertainty and Ambiguity”, *Journal of Risk and Uncertainty*, **5**, 325–370.
- CHARNESS, G. and LEVIN, D. (2005), “When Optimal Choices Feel Wrong: A Laboratory Study of Bayesian Updating, Complexity, and Affect”, *American Economic Review*, **95**, 1300–1309.
- CHEW, S. H. and TAN, G. (2005), “The Market for Sweepstakes”, *Review of Economic Studies*, **72**, 1009–1029.
- CLEMEN, R. T. and LICHTENDAHL, K. C. (2005), *Debiasing Expert Overconfidence: A Bayesian Calibration Model* (Durham, NC: Fuqua School of Business, Duke University).
- CLEMEN, R. T. and ROLLE, F. (2001), “In Theory ... in Practice”, *Decision Analysis Newsletter*, **20** (1), 3.
- DE FINETTI, B. (1937), “La Prévision: Ses Lois Logiques, ses Sources Subjectives”, *Annales de l’Institut Henri Poincaré*, **7**, 1–68.
- ECHTERNACHT, G. J. (1972), “The Use of Confidence Testing in Objective Tests”, *Review of Educational Research*, **42**, 217–236.
- EDWARDS, W. (1954), “The Theory of Decision Making”, *Psychological Bulletin*, **51**, 380–417.
- ELLSBERG, D. (1961), “Risk, Ambiguity and the Savage Axioms”, *Quarterly Journal of Economics*, **75**, 643–669.
- GHIRARDATO, P. and MARINACCI, M. (2001), “Risk, Ambiguity, and the Separation of Utility and Beliefs”, *Mathematics of Operations Research*, **26**, 864–890.
- GILBOA, I. (1987), “Expected Utility with Purely Subjective Non-Additive Probabilities”, *Journal of Mathematical Economics*, **16**, 65–88.
- GILBOA, I. (ed.) (2004), *Uncertainty in Economic Theory: Essays in Honor of David Schmeidler’s 65th Birthday* (London: Routledge).
- GILBOA, I. and SCHMEIDLER, D. (1989), “Maxmin Expected Utility with a Non-Unique Prior”, *Journal of Mathematical Economics*, **18**, 141–153.
- GOEREE, J. K., HOLT, C. A. and PALFREY, T. R. (2002), “Quantal Response Equilibrium and Overbidding in Private-Value Auctions”, *Journal of Economic Theory*, **104**, 247–272.
- GONZALEZ, R. and WU, G. (1999), “On the Shape of the Probability Weighting Function”, *Cognitive Psychology*, **38**, 129–166.

- GONZALEZ, R. and WU, G. (2003), "Composition Rules in Original and Cumulative Prospect Theory" (Mimeo).
- GOOD, I. J. (1952), "Rational Decisions", *Journal of the Royal Statistical Society Series B*, **14**, 107–114.
- GREENSPAN, A. (2004), "Innovations and Issues in Monetary Policy: The Last Fifteen Years", *American Economic Review, Papers and Proceedings*, **94**, 33–40.
- GUL, F. (1991), "A Theory of Disappointment Aversion", *Econometrica*, **59**, 667–686.
- HALEVY, Y. and FELTKAMP, V. (2005), "A Bayesian Approach to Uncertainty Aversion", *Review of Economic Studies*, **72**, 449–466.
- HALEVY, Y. and OZDENOREN, E. (2007), "Uncertainty and Compound Lotteries: Calibration" (Working Paper, University of British Columbia).
- HANSON, R. (2002), "Piece Entitled 'Wanna bet?'" *Nature*, **420** (November), 354–355.
- HARRISON, G. W., LAU, M. I. and WILLIAMS, M. B. (2002), "Estimating Individual Discount Rates in Denmark: A Field Experiment", *American Economic Review*, **92**, 1606–1617.
- HOLT, C. A. (1986), "Preference Reversals and the Independence Axiom", *American Economic Review*, **76**, 508–513.
- HOLT, C. A. (2006), *Webgames and Strategy: Recipes for Interactive Learning* (London: Addison-Wesley).
- HOLT, C. A. and LAURY, S. K. (2002), "Risk Aversion and Incentive Effects", *American Economic Review*, **92**, 1644–1655.
- HUCK, S. and WEIZSÄCKER, G. (2002), "Do Players Correctly Estimate What Others Do? Evidence of Conservatism in Beliefs", *Journal of Economic Behavior and Organization*, **47**, 71–85.
- HURWICZ, L. (1960), "Optimality and Informational Efficiency in Resource Allocation", in Arrow, K. J., Karlin, S. and Suppes, P. (eds.) *Mathematical Methods in the Social Sciences* (Stanford, CA: Stanford University Press) 17–46.
- JOHNSTONE, D. J. (2007a), "The Value of Probability Forecast from Portfolio Theory", *Theory and Decision*, **63**, 153–203.
- JOHNSTONE, D. J. (2007b), "Economic Darwinism: Who has the Best Probabilities", *Theory and Decision*, **62**, 47–96.
- JOUINI, E. and NAPP, C. (2007), "Consensus Consumer and Intertemporal Asset Pricing with Heterogeneous Beliefs", *Review of Economic Studies*, **74**, 1149–1174.
- KAHNEMAN, D. and TVERSKY, A. (1979), "Prospect Theory: An Analysis of Decision Under Risk", *Econometrica*, **47**, 263–291.
- KARNI, E. and SAFRA, Z. (1987), "Preference Reversal and the Observability of Preferences by Experimental Methods", *Econometrica*, **55**, 675–685.
- KEREN, G. B. (1991), "Calibration and Probability Judgments: Conceptual and Methodological Issues", *Acta Psychologica*, **77**, 217–273.
- KEYNES, J. M. (1921), *A Treatise on Probability* (London: Macmillan).
- KLIBANOFF, P., MARINACCI, M. and MUKERJI, S. (2005), "A Smooth Model of Decision Making Under Ambiguity", *Econometrica*, **73**, 1849–1892.
- KNIGHT, F. H. (1921), *Risk, Uncertainty, and Profit* (New York: Houghton Mifflin).
- LEE, J. (2008), "The Effect of the Background Risk in a Simple Chance Improving Decision Model", *The Journal of Risk and Uncertainty*, **36**, 19–41.
- LI, W. (2007), "Changing One's Mind When the Facts Change: Incentives of Experts and the Design of Reporting Protocols", *Review of Economic Studies*, **74**, 1175–1194.
- LUCE, R. D. (1959), *Individual Choice Behavior* (New York: Wiley).
- LUCE, R. D. (1991), "Rank- and Sign-Dependent Linear Utility Models for Binary Gambles", *Journal of Economic Theory*, **53**, 75–100.
- LUCE, R. D. (2000), *Utility of Gains and Losses: Measurement-Theoretical and Experimental Approaches* (London: Lawrence Erlbaum).
- MACCHERONI, F., MARINACCI, M. and RUSTICHINI, A. (2006), "Ambiguity Aversion, Robustness, and the Variational Representation of Preferences", *Econometrica*, **74**, 1447–1498.
- MACHINA, M. J. (2004), "Almost-Objective Uncertainty", *Economic Theory*, **24**, 1–54.
- MACHINA, M. J. and SCHMEIDLER, D. (1992), "A More Robust Definition of Subjective Probability", *Econometrica*, **60**, 745–780.
- McFADDEN, D. L. (1974), "Conditional Logit Analysis of Qualitative Choice Behavior", in Zarembka, P. (ed.) *Frontiers of Econometrics* (New York: Academic Press) 105–142.
- McFADDEN, D. L. (1976), "Quantal Choice Analysis: A Survey", *Annals of Economic and Social Measurement*, **5**, 363–390.
- MANSKI, C. F. (2004), "Measuring Expectations", *Econometrica*, **72**, 1329–1376.
- McCLELLAND, A. and BOLGER, F. (1994), "The Calibration of Subjective Probabilities: Theories and Models 1980–1994", in Wright, G. and Ayton, P. (eds.) *Subjective Probability* (New York: Wiley) 453–481.
- McKELVEY, R. and PAGE, T. (1990), "Public and Private Information: An Experimental Study of Information Pooling", *Econometrica*, **58**, 1321–1339.
- MIYAMOTO, J. M. (1988), "Generic Utility Theory: Measurement Foundations and Applications in Multiattribute Utility Theory", *Journal of Mathematical Psychology*, **32**, 357–404.
- MOSTELLER, F. and NOGEE, P. (1951), "An Experimental Measurement of Utility", *Journal of Political Economy*, **59**, 371–404.

- MURPHY, A. H. and WINKLER, R. L. (1974), "Subjective Probability Forecasting Experiments in Meteorology: Some Preliminary Results", *Bulletin of the American Meteorological Society*, **55**, 1206–1216.
- MYAGKOV, M. G. and PLOTT, C. R. (1997), "Exchange Economies and Loss Exposure: Experiments Exploring Prospect Theory and Competitive Equilibria in Market Environments", *American Economic Review*, **87**, 801–828.
- NAU, R. F. (2006), "Uncertainty Aversion with Second-Order Utilities and Probabilities", *Management Science*, **52**, 136–145.
- NYARKO, Y. and SCHOTTER, A. (2002), "An Experimental Study of Belief Learning Using Elicited Beliefs", *Econometrica*, **70**, 971–1005.
- OLSZEWSKI, W. (2007), "Preferences over Sets of Lotteries", *Review of Economic Studies*, **74**, 567–595.
- PALFREY, T. R. and WANG, S. W. (2007), "On Eliciting Beliefs in Strategic Games. Division of the Humanities and Social Sciences", CalTech, Pasadena, CA 91125.
- PALMER, T. N. and HAGEDORN, R. (eds.) (2006), *Predictability of Weather and Climate* (Cambridge: Cambridge University Press).
- PFANZAGL, J. (1959), "A General Theory of Measurement—Applications to Utility", *Naval Research Logistics Quarterly*, **6**, 283–294.
- PRELEC, D. (1998), "The Probability Weighting Function", *Econometrica*, **66**, 497–527.
- PRELEC, D. (2004), "A Bayesian Truth Serum for Subjective Data", *Science*, **306**, 462–466.
- QUIGGIN, J. (1982), "A Theory of Anticipated Utility", *Journal of Economic Behaviour and Organization*, **3**, 323–343.
- RAIFFA, H. (1968), *Decision Analysis* (London: Addison-Wesley).
- SANDRONI, A., SMORODINSKY, R. and VOHRA, R. V. (2003), "Calibration with Many Checking Rules", *Mathematics of Operations Research*, **28**, 141–153.
- SAVAGE, L. J. (1954), *The Foundations of Statistics* (New York: Wiley).
- SAVAGE, L. J. (1971), "Elicitation of Personal Probabilities and Expectations", *Journal of the American Statistical Association*, **66**, 783–801.
- SCHMEIDLER, D. (1989), "Subjective Probability and Expected Utility without Additivity", *Econometrica*, **57**, 571–587.
- SEGAL, U. and SPIVAK, A. (1990), "First-Order Versus Second-Order Risk-Aversion", *Journal of Economic Theory*, **51**, 111–125.
- SELTEN, R., SADRIEH, A. and ABBINK, K. (1999), "Money does not Induce Risk Neutral Behavior, but Binary Lotteries do Even Worse", *Theory and Decision*, **46**, 211–249.
- SHAFER, G. (1976), *A Mathematical Theory of Evidence* (NJ: Princeton University Press).
- SHILLER, R. J., KON-YA, F. and TSUTSUI, Y. (1996), "Why did the Nikkei Crash? Expanding the Scope of Expectations Data Collection", *The Review of Economics and Statistics*, **78**, 156–164.
- SPIEGELHALTER, D. J. (1986), "Probabilistic Prediction in Patient Management and Clinical Trials", *Statistics in Medicine*, **5**, 421–433.
- STAEL VON HOLSTEIN, C. A. S. (1972), "Probabilistic Forecasting: An Experiment Related to the Stock Market", *Organizational Behaviour and Human Performance*, **8**, 139–158.
- STARMER, C. and SUGDEN, R. (1991), "Does the Random-Lottery Incentive System Elicit True Preferences? An Experimental Investigation", *American Economic Review*, **81**, 971–978.
- TETLOCK, P. E. (2005), *Expert Political Judgment* (NJ: Princeton University Press).
- THALER, R. H. and JOHNSON, E. J. (1990), "Gambling with the House Money and Trying to Break Even: The Effects of Prior Outcomes on Risky Choice", *Management Science*, **36**, 643–660.
- TVERSKY, A. and KAHNEMAN, D. (1992), "Advances in Prospect Theory: Cumulative Representation of Uncertainty", *Journal of Risk and Uncertainty*, **5**, 297–323.
- TVERSKY, A. and KOEHLER, D. J. (1994), "Support Theory: A Nonextensional Representation of Subjective Probability", *Psychological Review*, **101**, 547–567.
- VAN DE KUILEN, G. and WAKKER, P. P. (2009), *The Midweight Method to Measure Attitudes Towards Risk and Ambiguity* (Rotterdam, The Netherlands: Econometric Institute, Erasmus University).
- VON NEUMANN, J. and MORGENTERN, O. (1944, 1947, 1953), *Theory of Games and Economic Behavior* (Princeton, NJ: Princeton University Press).
- WAKKER, P. P. (2004), "On the Composition of Risk Preference and Belief", *Psychological Review*, **111**, 236–241.
- WAKKER, P. P. (2009), *Prospect Theory for Risk and Ambiguity* (Cambridge, UK: Cambridge University Press).
- WAKKER, P. P. and DENEFFE, D. (1996), "Eliciting von Neumann–Morgenstern Utilities When Probabilities are Distorted or Unknown", *Management Science*, **42**, 1131–1150.
- WINKLER, R. L. and MURPHY, A. H. (1970), "Nonlinear Utility and the Probability Score", *Journal of Applied Meteorology*, **9**, 143–148.
- WOLFERS, J. and ZITZEWITZ, E. (2004), "Prediction Markets", *Journal of Economic Perspectives*, **18**, 107–126.
- WRIGHT, W. F. (1988), "Empirical Comparison of Subjective Probability Elicitation Methods", *Contemporary Accounting*, **5**, 47–57.
- YATES, J. F. (1990), *Judgment and Decision Making* (London: Prentice Hall).